

DATA 598: Final Project Report

Hriday Baghar

June 8, 2022

Introduction

Recent news coverage of the World Air Quality Report by IQAir (a Swiss organization) showed that 35 Indian cities were in the top 50 of the world's most polluted cities (The Indian Express, 2022). It has been established and extensively documented that exposure to high amounts of air pollution leads to serious respiratory problems. Records show that in 2019 over 1.6 million deaths were attributed to poor air quality. The cause of death ranged from strokes, diabetes, lung cancer and myocardial infarctions (IQAir, 2022).

Poor air quality has become a cause for concern in recent years where the city of Delhi, India had to enforce stay at home guidelines (Hindustan Times, 2021) and shut down schools (DW, 2021). It is clear that there is a need to develop solutions and take measures to bring air pollution under control, so that people can live their lives without disruption and without the risk of illness.

Models for the series can help identify:

- seasonal patterns of different frequency in pollution and air quality and serve as a starting point for policymakers to investigate sources of air pollution and how they can be controlled
- trends of air quality to check how interventions have impacted air quality in the short and long term

```
library(fpp3)
library(astsa) # for muspec()

data <- read.csv("data/city_day_raw.csv") |>
  select(c(City, Date, AQI)) |>
  mutate(Date = as.Date(Date, format="%Y-%m-%d")) |>
  filter(City %in% c("Bengaluru", "Chennai", "Delhi", "Hyderabad", "Lucknow"))
```

Exploratory Data Analysis

a. Plot the series

```
ggplot(data = data, aes(x=Date, y=AQI, color=City)) +
  geom_line()+
  facet_grid(City ~.)
```

```
## Warning: Removed 326 row(s) containing missing values (geom_path).
```



b. Series description

Observations:

- Some series have missing values in the start (we will remove them)
- Some series also have missing values in between - we will impute them with the previous value (downward filling)
- For the first two years, Hyderabad data has very big outliers compared to the rest of the years (possible that something about the sensors changed)
- Seasonality
 - Clear annual seasonal pattern in Delhi and Lucknow - cities in northern India
 - Not so clear in Chennai, Bengaluru and Hyderabad - cities in southern India
 - Might be worth looking into some domain knowledge regarding air quality

Handling missing values

```
data |>
  group_by(City) |>
  summarise(missing.percent = sum(is.na(AQI))/n()*100)
```

```
## # A tibble: 5 x 2
##   City      missing.percent
##   <chr>      <dbl>
## 1 Bengaluru      4.93
## 2 Chennai        6.22
## 3 Delhi          0.498
## 4 Hyderabad      6.28
```

```
## 5 Lucknow          5.77
```

Most of the missing values are likely due to the long runs of NAs in the beginning for each city. We can fill the NAs in between with the previous value and filter out the starting runs of NAs.

First, let us fill “downwards” i.e. only for dates which have data on the previous day. This will fill the previous day’s value in the missing field.

```
data |>
  group_by(City) |>
  fill(AQI, .direction = "down") |>
  summarise(missing.percent = sum(is.na(AQI))/n()*100)
```

```
## # A tibble: 5 x 2
##   City      missing.percent
##   <chr>         <dbl>
## 1 Bengaluru      3.93
## 2 Chennai       4.08
## 3 Delhi          0
## 4 Hyderabad     4.29
## 5 Lucknow       3.93
```

We now see that most of the missing data is only from the long runs of NA in the start. We can safely filter those records out.

We clean the missing values in the following manner:

- Group data by City
- Fill missing values downwards
- Remove remaining missing values at the start

```
(data <- data |>
  group_by(City) |>
  fill(AQI, .direction = "down") |>
  filter(!is.na(AQI)))
```

```
## # A tibble: 9,716 x 3
## # Groups:   City [5]
##   City      Date      AQI
##   <chr>    <date>    <dbl>
## 1 Bengaluru 2015-03-21    91
## 2 Bengaluru 2015-03-22   120
## 3 Bengaluru 2015-03-23   154
## 4 Bengaluru 2015-03-24   119
## 5 Bengaluru 2015-03-25   232
## 6 Bengaluru 2015-03-26   132
## 7 Bengaluru 2015-03-27   123
## 8 Bengaluru 2015-03-28   152
## 9 Bengaluru 2015-03-29   143
## 10 Bengaluru 2015-03-30    80
## # ... with 9,706 more rows
```

```
ggplot(data = data, aes(x=Date, y=AQI, color=City)) +
  geom_line()+
  facet_grid(City ~.)
```



c. Evaluating Stationarity

```
data <- data |>
  as_tsibble(index = Date, key = City)

data |>
  features(AQI, unitroot_ndiffs)
```

```
## # A tibble: 5 x 2
##   City      ndiffs
##   <chr>      <int>
## 1 Bengaluru     1
## 2 Chennai      1
## 3 Delhi        1
## 4 Hyderabad    1
## 5 Lucknow      1
```


We use the `ndiffs` function to evaluate stationarity. It is clear that each series is non-stationary and would require one differencing operation. We will require to perform this again when we decide to modify the level of our data.

d. Investigating Seasonality

We use spectral analysis to identify if there are multiple seasonal patterns. Based on our observations we will decide how to roll up the data.


```
delhi.spec <- mvspec(data[data$City == "Delhi", "AQI"], detrend = TRUE, spans = 5)
```

Series: data[data\$City == "Delhi", "AQI"] | Smoothed Periodogram | taper




```
lucknow.spec <- mvspec(data[data$City == "Lucknow", "AQI"], detrend = TRUE, spans = 5)
```

series: data[data\$City == "Lucknow", "AQI"] | Smoothed Periodogram | tap




```
chennai.spec <- mvspec(data[data$City == "Chennai", "AQI"], detrend = TRUE, spans = 5)
```

Series: data[data\$City == "Chennai", "AQI"] | Smoothed Periodogram | tap




```
bengaluru.spec <- mvspec(data[data$City == "Bengaluru", "AQI"], detrend = TRUE, spans = 5)
```

series: data[data\$City == "Bengaluru", "AQI"] | Smoothed Periodogram | tai



```
hyderabad.spec <- mvspec(data[data$City == "Hyderabad", "AQI"], detrend = TRUE, spans = 5)
```


series: data[data\$City == "Hyderabad", "AQI"] | Smoothed Periodogram | ta



Initial plots suggest that Delhi and Lucknow have a single seasonal pattern, probably the strong annual pattern that we saw in the plots. Hyderabad is also suggestive of a single seasonal pattern. We could see a weak pattern in the plots so this should not be surprising.

Bengaluru and Chennai have slightly more complicated plots

Let us observe all peaks. Peak thresholds have been listed for each series as comments based on the spectral plots above

```
get.spec.df <- function(spec){
  return(data.frame(spec$details))
}

delhi.spec <- get.spec.df(delhi.spec)
bengaluru.spec <- get.spec.df(bengaluru.spec)
chennai.spec <- get.spec.df(chennai.spec)
lucknow.spec <- get.spec.df(lucknow.spec)
hyderabad.spec <- get.spec.df(hyderabad.spec)

# Anything > 400,000 for Delhi
delhi.spec |> filter(spectrum > 400000) |> arrange(desc(spectrum))
```

```
## frequency period spectrum
## 1 0.0030 337.5000 1429770.1
## 2 0.0025 405.0000 1405032.7
## 3 0.0035 289.2857 1192047.7
## 4 0.0020 506.2500 974836.9
## 5 0.0054 184.0909 611925.1
## 6 0.0049 202.5000 609912.1
```

```
## 7    0.0059 168.7500 586935.5
## 8    0.0040 253.1250 567373.7
## 9    0.0044 225.0000 401246.1
```

Anything > 250,000 for Lucknow

```
lucknow.spec |> filter(spectrum > 40000) |> arrange(desc(spectrum))
```

```
##      frequency    period  spectrum
## 1      0.0026   388.8000 1639476.58
## 2      0.0031   324.0000 1635788.29
## 3      0.0021   486.0000 1465784.65
## 4      0.0036   277.7143 1065902.83
## 5      0.0015   648.0000  777092.90
## 6      0.0051   194.4000  463143.71
## 7      0.0057   176.7273  424077.73
## 8      0.0041   243.0000  397417.85
## 9      0.0046   216.0000  391343.55
## 10     0.0005  1944.0000  339813.12
## 11     0.0062   162.0000  335642.74
## 12     0.0010   972.0000  279736.97
## 13     0.0067   149.5385  141288.25
## 14     0.0113    88.3636   60118.21
## 15     0.0118    84.5217   56434.30
## 16     0.0108    92.5714   54759.99
## 17     0.0103    97.2000   43767.60
```

Anything > 15,000 for Chennai

```
chennai.spec |> filter(spectrum > 15000) |> arrange(desc(spectrum))
```

```
##      frequency    period spectrum
## 1      0.0031   324.0000 47886.86
## 2      0.0036   277.7143 46776.24
## 3      0.0026   388.8000 44059.75
## 4      0.0041   243.0000 38142.38
## 5      0.0005  1944.0000 36850.41
## 6      0.0010   972.0000 30704.17
## 7      0.0021   486.0000 28071.32
## 8      0.0046   216.0000 24100.55
## 9      0.0015   648.0000 23335.65
## 10     0.0051   194.4000 19271.26
## 11     0.0057   176.7273 18727.97
## 12     0.0283    35.3455 18541.12
## 13     0.0195    51.1579 18209.25
## 14     0.0278    36.0000 18052.08
## 15     0.0190    52.5405 17930.79
## 16     0.0273    36.6792 17575.54
## 17     0.0226    44.1818 17340.32
## 18     0.0221    45.2093 16655.65
## 19     0.0231    43.2000 16547.35
## 20     0.0062   162.0000 16478.47
## 21     0.0201    49.8462 15037.00
## 22     0.0267    37.3846 15026.31
```

Anything > 20,000 for Bengaluru

```
bengaluru.spec |> filter(spectrum > 20000) |> arrange(desc(spectrum))
```

```
##      frequency    period spectrum
## 1      0.0031   324.0000 55841.12
## 2      0.0036   277.7143 47777.66
## 3      0.0026   388.8000 45733.27
## 4      0.0021   486.0000 36283.75
## 5      0.0041   243.0000 34004.44
## 6      0.0005  1944.0000 26091.02
## 7      0.0015   648.0000 25401.20
## 8      0.0046   216.0000 23792.16
## 9      0.0010   972.0000 20947.71
```

```
# Anything > 25,000 for Hyderabad
hyderabad.spec |> filter(spectrum > 25000) |> arrange(desc(spectrum))
```

```
##      frequency    period  spectrum
## 1      0.0031  320.0000 122896.69
## 2      0.0026  384.0000 113574.83
## 3      0.0036  274.2857  96186.72
## 4      0.0021  480.0000  96000.89
## 5      0.0042  240.0000  49442.58
## 6      0.0016  640.0000  43115.11
## 7      0.0047  213.3333  35542.95
## 8      0.0052  192.0000  27722.66
## 9      0.0073  137.1429  27491.60
## 10     0.0068  147.6923  26661.93
```

Observations:

- Delhi, Lucknow and Hyderabad appear to have a seasonal pattern that occurs roughly annually (all 3 have peaks around 340 - 400 days)
- Bengaluru and Chennai exhibit relatively much weaker seasonal behavior but this seasonal pattern also appears to occur annually.

Given the above observations, we can model the seasonality as annual. For ARIMA modeling, we should roll up our data to a monthly level and model the seasonal period as 12.

```
data <- data |>
  as_tibble() |>
  mutate(Date = yearmonth(Date)) |>
  group_by(City, Date) |>
  summarise(AQI = mean(AQI)) |>
  as_tsibble(index = Date, key = City)
```


```
## `summarise()` has grouped output by 'City'. You can override using the
## `.groups` argument.
```

```
write.csv(x = data |> as_tibble(), file = "data/city_month_cleaned.csv")
```

Plotting rolled up data:

```
data |> autoplot(AQI) + facet_grid(City ~ .)
```


```
## `mutate_if()` ignored the following grouping variables:
## * Column `City`
```



Re-analyzing seasonal pattern on rolled up data


```
delhi.spec <- mvspec(data[data$City == "Delhi", "AQI"], detrend = TRUE, spans = 2)
```

Series: data[data\$City == "Delhi", "AQI"] | Smoothed Periodogram | taper




```
lucknow.spec <- mvspec(data[data$City == "Lucknow", "AQI"], detrend = TRUE, spans = 2)
```

series: data[data\$City == "Lucknow", "AQI"] | Smoothed Periodogram | tap




```
chennai.spec <- mvspec(data[data$City == "Chennai", "AQI"], detrend = TRUE, spans = 2)
```

Series: data[data\$City == "Chennai", "AQI"] | Smoothed Periodogram | tap




```
bengaluru.spec <- mvspec(data[data$City == "Bengaluru", "AQI"], detrend = TRUE, spans = 2)
```

series: data[data\$City == "Bengaluru", "AQI"] | Smoothed Periodogram | tai



```
hyderabad.spec <- mvspec(data[data$City == "Hyderabad", "AQI"], detrend = TRUE, spans = 2)
```


series: data[data\$City == "Hyderabad", "AQI"] | Smoothed Periodogram | ta



```
delhi.spec <- get.spec.df(delhi.spec)
bengaluru.spec <- get.spec.df(bengaluru.spec)
chennai.spec <- get.spec.df(chennai.spec)
lucknow.spec <- get.spec.df(lucknow.spec)
hyderabad.spec <- get.spec.df(hyderabad.spec)
```

```
head(delhi.spec |> arrange(desc(spectrum)))
```

```
##   frequency  period spectrum
## 1    0.0833 12.0000 92039.05
## 2    0.0972 10.2857 49442.67
## 3    0.0694 14.4000 48518.62
## 4    0.1667  6.0000 31787.03
## 5    0.1528  6.5455 21819.76
## 6    0.1806  5.5385 15642.56
```

```
head(lucknow.spec |> arrange(desc(spectrum)))
```

```
##   frequency  period spectrum
## 1    0.0833 12.0000 118009.11
## 2    0.0694 14.4000  63179.60
## 3    0.0972 10.2857  60057.08
## 4    0.1667  6.0000  21361.00
## 5    0.0139 72.0000  13362.27
## 6    0.1806  5.5385  12966.35
```

```
head(chennai.spec |> arrange(desc(spectrum)))
```

```
##   frequency  period spectrum
```

```
## 1    0.0139 72.0000 3273.838
## 2    0.1111  9.0000 2879.934
## 3    0.0972 10.2857 2640.238
## 4    0.0833 12.0000 1911.123
## 5    0.0556 18.0000 1892.193
## 6    0.0278 36.0000 1822.395
```

```
head(bengaluru.spec |> arrange(desc(spectrum)))
```

```
##   frequency  period  spectrum
## 1    0.0833 12.0000 2414.1561
## 2    0.0972 10.2857 1698.7839
## 3    0.0694 14.4000 1453.2423
## 4    0.1111  9.0000 1307.5155
## 5    0.1250  8.0000 1035.6042
## 6    0.0278 36.0000  964.7626
```

```
head(hyderabad.spec |> arrange(desc(spectrum)))
```

```
##   frequency  period  spectrum
## 1    0.0833 12.0000 6619.792
## 2    0.0972 10.2857 4845.133
## 3    0.0694 14.4000 3318.923
## 4    0.1111  9.0000 1940.089
## 5    0.2222  4.5000 1484.920
## 6    0.1250  8.0000 1439.564
```

Our hypothesis seems to be correct. We see a clear 12 month seasonal pattern in all cities except Chennai. We also see a smaller pattern in all series other than Chennai, but those peaks are much smaller than the 12 month peaks. For now we choose not to include them in our seasonal pattern.

The 72 month pattern for Chennai does not make sense as the data itself covers 5 years. Chennai does seem to exhibit seasonal patterns of 9, 10.28, 12 and 18 months. Contrary to other cities, the seasonal pattern does not have one clear influence. It is worth noting that Chennai is the only coastal city out of all the cities in the data. Perhaps this influences the seasonality of AQI? More investigation and comparison of other coastal cities (specifically on India's east coast) would be required to make a clear conclusion.

e. ACF and PACF

```
# Re-evaluating stationarity
```

```
data |> features(AQI, unitroot_kpss)
```

```
## # A tibble: 5 x 3
##   City      kpss_stat kpss_pvalue
##   <chr>      <dbl>      <dbl>
## 1 Bengaluru  0.630      0.0199
## 2 Chennai   1.04       0.01
## 3 Delhi     0.387     0.0828
## 4 Hyderabad 0.866     0.01
## 5 Lucknow   0.115     0.1
```

```
data |> features(AQI, unitroot_ndiffs)
```

```
## # A tibble: 5 x 2
##   City      ndiffs
##   <chr>      <int>
```

```
## 1 Bengaluru      1
## 2 Chennai        1
## 3 Delhi          0
## 4 Hyderabad      1
## 5 Lucknow        0

data |> features(AQI, unitroot_nsdiffs)
```

```
## # A tibble: 5 x 2
##   City      nsdiffs
##   <chr>      <int>
## 1 Bengaluru    0
## 2 Chennai     0
## 3 Delhi       1
## 4 Hyderabad   0
## 5 Lucknow     1
```

We see that Delhi and Lucknow require a seasonal difference and the others require a non-seasonal difference.

```
data |>
  filter(City %in% c("Delhi", "Lucknow")) |>
  features(difference(AQI, lag = 12), unitroot_ndiffs)
```


```
## # A tibble: 2 x 2
##   City      ndiffs
##   <chr>      <int>
## 1 Delhi       0
## 2 Lucknow     1
```

Lucknow also requires a non-seasonal difference after a seasonal difference.

```
data |> filter(City == "Delhi") |> gg_tsdisplay(difference(AQI, lag = 12),
  plot_type = "partial")
```

```
## Warning: Removed 12 row(s) containing missing values (geom_path).
```


```
## Warning: Removed 12 rows containing missing values (geom_point).
```



```
data |> filter(City == "Lucknow") |> gg_tsdisplay(difference(
  difference(AQI, lag=12),
  lag = 1),
  plot_type = "partial")
```

```
## Warning: Removed 13 row(s) containing missing values (geom_path).
```




```
## Warning: Removed 13 rows containing missing values (geom_point).
```



```
data |> filter(City == "Chennai") |> gg_tsdisplay(difference(AQI),
  plot_type = "partial")
```




```
## Warning: Removed 1 row(s) containing missing values (geom_path).
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```




```
data |> filter(City == "Bengaluru") |> gg_tsdisplay(difference(AQI),
                                                    plot_type = "partial")
```

```
## Warning: Removed 1 row(s) containing missing values (geom_path).
## Removed 1 rows containing missing values (geom_point).
```



```
data |> filter(City == "Hyderabad") |> gg_tsddisplay(difference(AQI),
  plot_type = "partial")
```

```
## Warning: Removed 1 row(s) containing missing values (geom_path).
## Removed 1 rows containing missing values (geom_point).
```



Candidate models:

- Delhi: $\text{pdq}(1,0,1) + \text{PDQ}(1,1,1)[12]$
 - Clear AR and MA spike at 1, seasonal AR and MA spike at 12
- Lucknow: $\text{pdq}(1,1,1) + \text{PDQ}(1,1,1)[12]$
 - Clear AR and MA spike at 1, seasonal AR and MA spike at 12
- Chennai: $\text{pdq}(0,1,0) + \text{PDQ}(0,0,0)[12]$
 - No significant spikes whatsoever
- Bengaluru: $\text{pdq}(0,1,0) + \text{PDQ}(0,0,1)[12]$
 - Seasonal MA spike at 12 in ACF
- Hyderabad: $\text{pdq}(0,1,0) + \text{PDQ}(0,0,0)[12]$
 - No significant spikes whatsoever
 - We will iterate with a seasonal term since we saw a weak seasonal pattern

ARIMA Modeling

Splitting data into train and test

```
train <- data |> filter(Date < yearmonth("2019 Jun"))
test <- data |> filter(Date >= yearmonth("2019 Jun"))
```

We now build the candidate models and display their fit metrics and residuals. It is clear from the plots below that all the series are white noise.

Delhi: $\text{pdq}(1,0,1) + \text{PDQ}(1,1,1)[12]$


```
delhi.fit <- train |>
  filter(City == "Delhi") |>
  model(arima = ARIMA(AQI ~ pdq(1,0,1) + PDQ(1,1,1)))

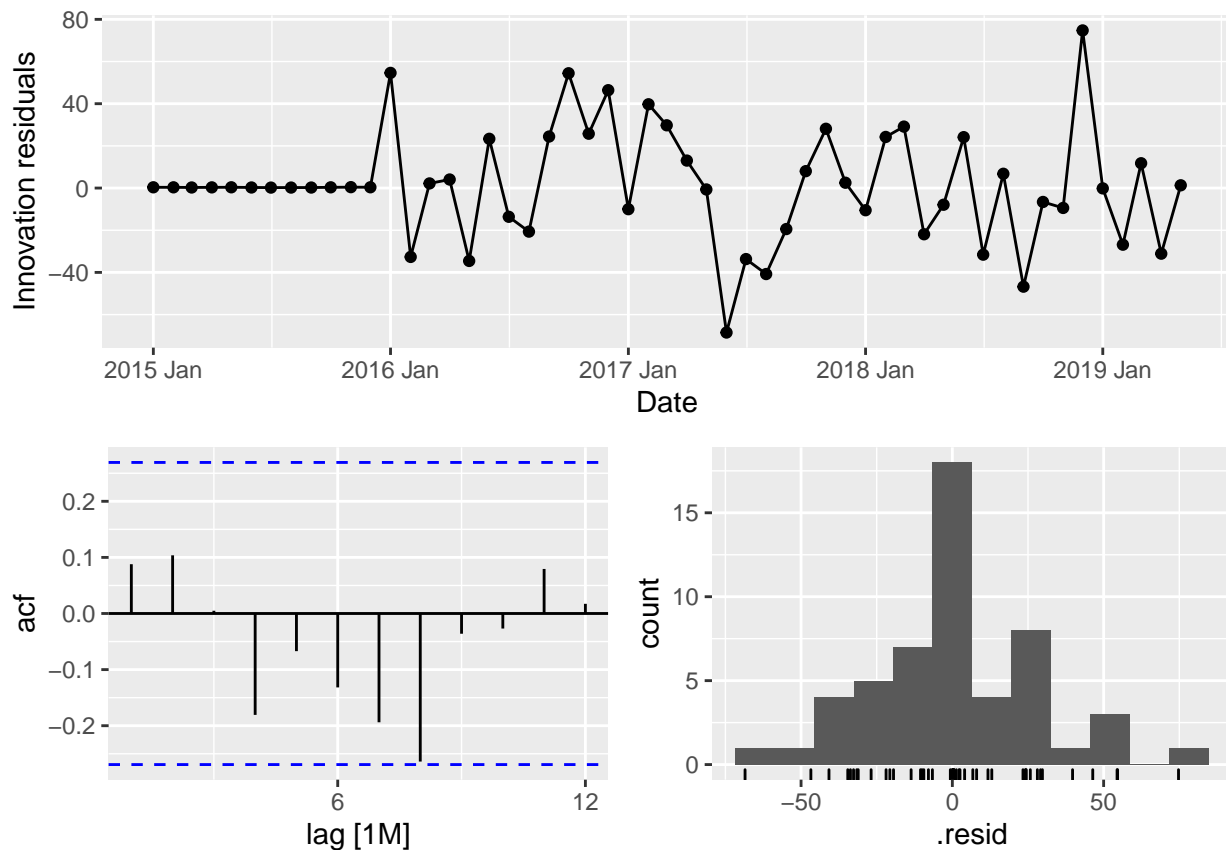
glance(delhi.fit)

## # A tibble: 1 x 9
##   City .model sigma2 log_lik   AIC   AICc   BIC ar_roots   ma_roots
##   <chr> <chr>   <dbl>   <dbl> <dbl> <dbl> <dbl> <list>   <list>
## 1 Delhi arima  1048.   -210.  431.  434.  441. <cpl [13]> <cpl [13]>

report(delhi.fit)

## Series: AQI
## Model: ARIMA(1,0,1)(1,1,1)[12] w/ drift
##
## Coefficients:
##          ar1      ma1      sar1      sma1  constant
##          0.4098  0.1929 -0.4944 -0.8443 -18.1421
## s.e.      0.2453  0.2689   0.2375   1.5218   4.8199
##
## sigma^2 estimated as 1048: log likelihood=-209.59
## AIC=431.18   AICc=433.65   BIC=441.46

delhi.fit |> gg_tsresiduals(lag = 12)
```



Lucknow: $\text{pdq}(1,1,1) + \text{PDQ}(1,1,1)[12]$

```
lucknow.fit <- train |>
  filter(City == "Lucknow") |>
  model(arima = ARIMA(AQI ~ pdq(1,1,1) + PDQ(1,1,1)))




glance(lucknow.fit)
```

```
## # A tibble: 1 x 9
##   City      .model sigma2 log_lik   AIC  AICc   BIC ar_roots  ma_roots
##   <chr>   <chr>   <dbl>   <dbl> <dbl> <dbl> <dbl> <list>    <list>
## 1 Lucknow arima  1384.   -192.  394.  396.  402. <cpl [13]> <cpl [13]>
```

```
report(lucknow.fit)
```

```
## Series: AQI
## Model: ARIMA(1,1,1)(1,1,1)[12]
##
## Coefficients:
##      ar1      ma1      sar1      sma1
##    0.2035 -0.8997  0.1987 -0.6631
## s.e. 0.2043  0.1192  0.4572  0.5464
##
## sigma^2 estimated as 1384: log likelihood=-191.89
## AIC=393.78  AICc=395.65  BIC=401.97
```

```
lucknow.fit |> gg_tsresiduals(lag = 12)
```



Chennai: $\text{pdq}(0,1,0)+\text{PDQ}(0,0,0)[12]$

```
chennai.fit <- train |>
  filter(City == "Chennai") |>
  model(arima = ARIMA(AQI ~ pdq(0,1,0) + PDQ(0,0,0)))
```


```
glance(chennai.fit)
```

```
## # A tibble: 1 x 9
##   City      .model sigma2 log_lik   AIC   AICc   BIC ar_roots  ma_roots
##   <chr>   <chr>   <dbl>   <dbl> <dbl> <dbl> <dbl> <list>   <list>
## 1 Chennai arima    1316.   -251.  503.  503.  505. <cpl [0]> <cpl [0]>
```

```
report(chennai.fit)
```

```
## Series: AQI
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 1316: log likelihood=-250.5
## AIC=503.01 AICc=503.09 BIC=504.92
```

```
chennai.fit |> gg_tsresiduals(lag=12)
```



Bengaluru: $\text{pdq}(0,1,0)+\text{PDQ}(0,0,1)[12]$

```
bengaluru.fit <- train |>
  filter(City == "Bengaluru") |>
  model(arima = ARIMA(AQI ~ pdq(0,1,0) + PDQ(0,0,1)))
```

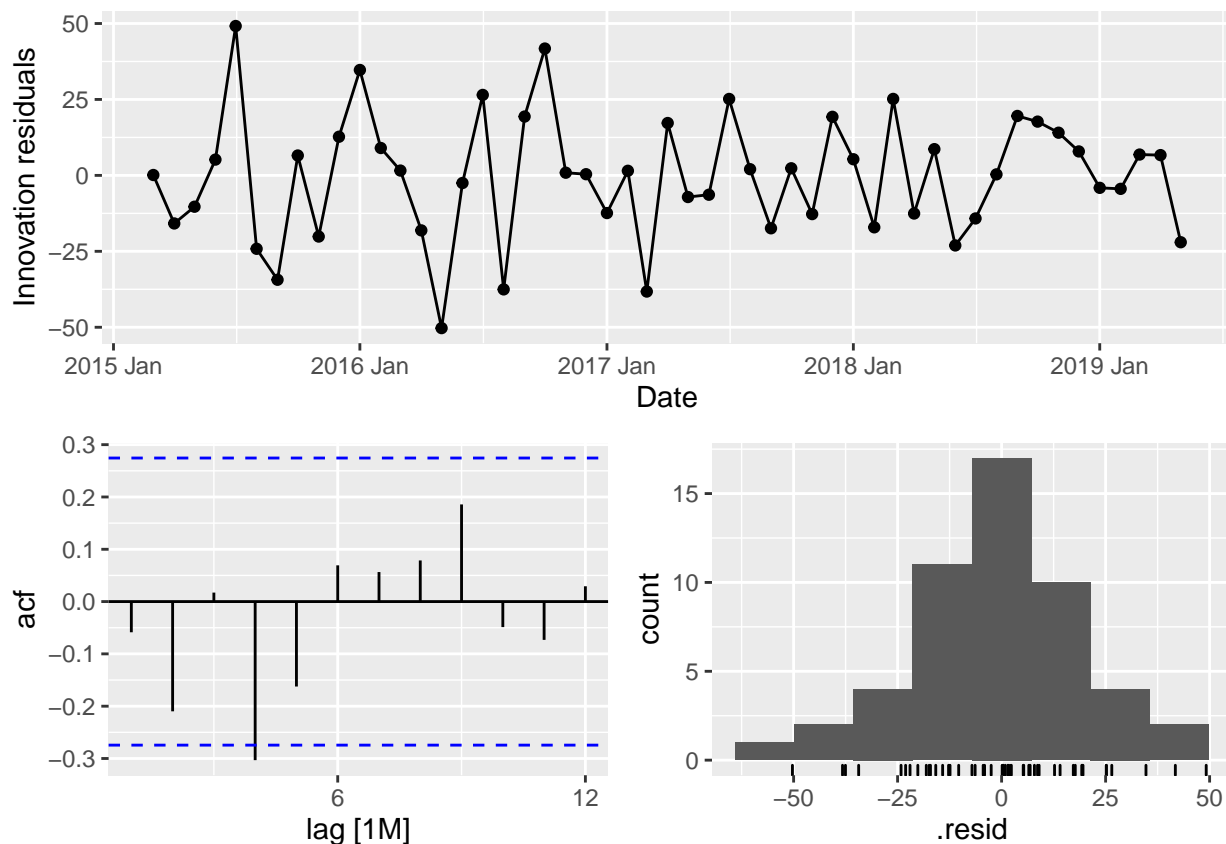
```
glance(bengaluru.fit)
```

```
## # A tibble: 1 x 9
##   City      .model sigma2 log_lik   AIC   AICc   BIC ar_roots  ma_roots
##   <chr>    <chr>   <dbl>  <dbl> <dbl> <dbl> <dbl> <list>   <list>
## 1 Bengaluru arima    419.   -222.  448.  449.  452. <cpl [0]> <cpl [12]>
```

```
report(bengaluru.fit)
```

```
## Series: AQI
## Model: ARIMA(0,1,0)(0,0,1)[12]
##
## Coefficients:
##      sma1
##      0.3586
## s.e. 0.1441
##
## sigma^2 estimated as 419.1: log likelihood=-222.22
## AIC=448.44 AICc=448.69 BIC=452.26
```

```
bengaluru.fit |> gg_tsresiduals(lag = 12)
```



Hyderabad: $\text{pdq}(0,1,0) + \text{PDQ}(0,0,0)[12]$

```
hyderabad.fit <- train |>
  filter(City == "Hyderabad") |>
  model(arima = ARIMA(AQI ~ pdq(0,1,0) + PDQ(0,0,0)))
```

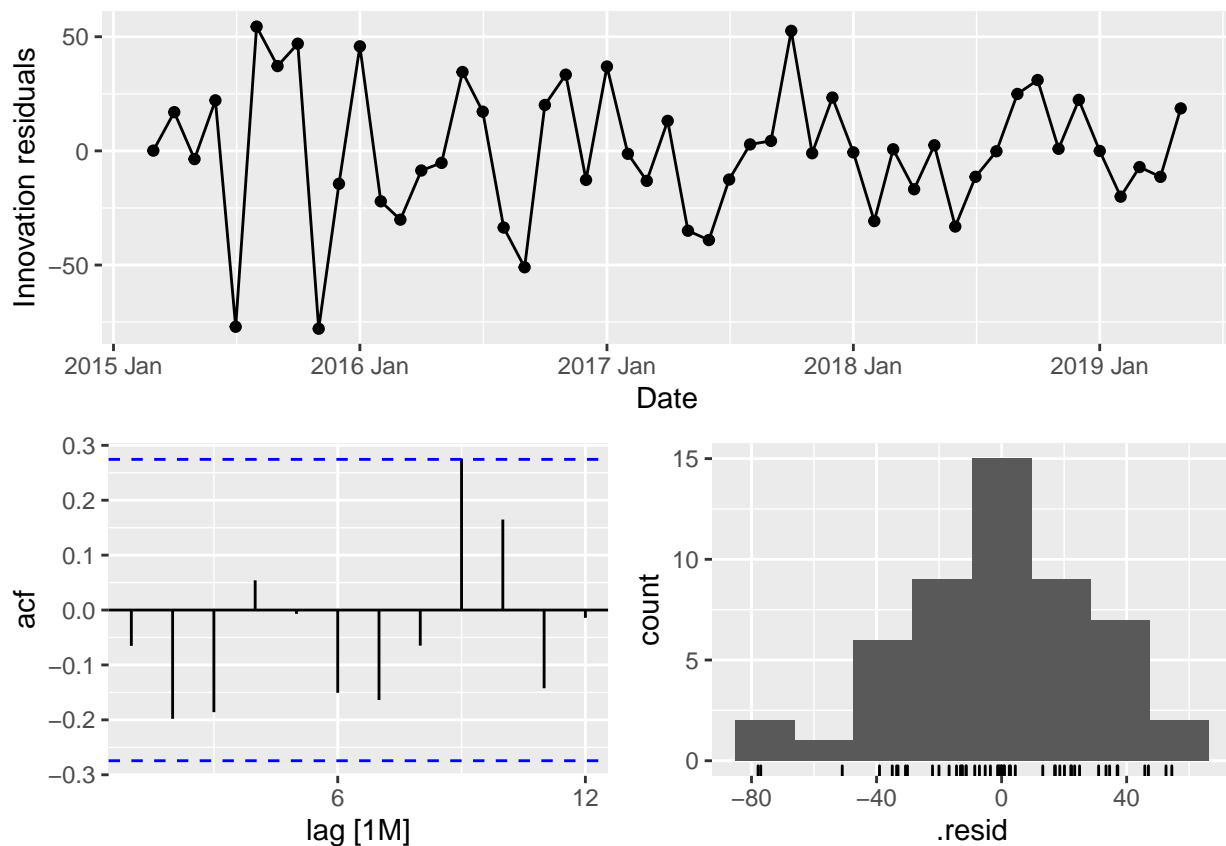
```
glance(hyderabad.fit)
```

```
## # A tibble: 1 x 9
##   City      .model sigma2 log_lik   AIC   AICc   BIC ar_roots  ma_roots
##   <chr>    <chr>   <dbl>  <dbl> <dbl> <dbl> <dbl> <list>   <list>
## 1 Hyderabad arima    870.  -240.  482.  482.  484. <cpl [0]> <cpl [0]>
```

```
report(hyderabad.fit)
```

```
## Series: AQI
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 869.7: log likelihood=-240.15
## AIC=482.3   AICc=482.38   BIC=484.21
```

```
hyderabad.fit |> gg_tsresiduals(lag=12)
```



Let us add a seasonal term in Hyderabad model since we saw that the series did exhibit weak seasonality. We compare models with 1 AR term only, 1 MA term only and 1 AR and MA term. Since we are including a seasonal term we use a seasonal difference on this model.

```
#MA term only
report(train |>
  filter(City == "Hyderabad") |>
  model(arima = ARIMA(AQI ~ pdq(0,0,0) + PDQ(0,1,1))))
```

```
## Series: AQI
## Model: ARIMA(0,0,0)(0,1,1)[12] w/ drift
##
```

```
## Coefficients:
##          sma1  constant
##        -0.7178 -13.6292
## s.e.    0.4353   3.2881
##
## sigma^2 estimated as 927.8:  log likelihood=-191.51
## AIC=389.02   AICc=389.71   BIC=394.01
```

```
#AR term only
report(train |>
  filter(City == "Hyderabad") |>
  model(arima = ARIMA(AQI ~ pdq(0,0,0) + PDQ(1,1,0))))
```

```
## Series: AQI
## Model: ARIMA(0,0,0)(1,1,0)[12] w/ drift
##
## Coefficients:
##          sar1  constant
##        -0.4717 -21.3408
## s.e.    0.1539   5.7460
##
## sigma^2 estimated as 1089:  log likelihood=-192.19
## AIC=390.39   AICc=391.07   BIC=395.38
```

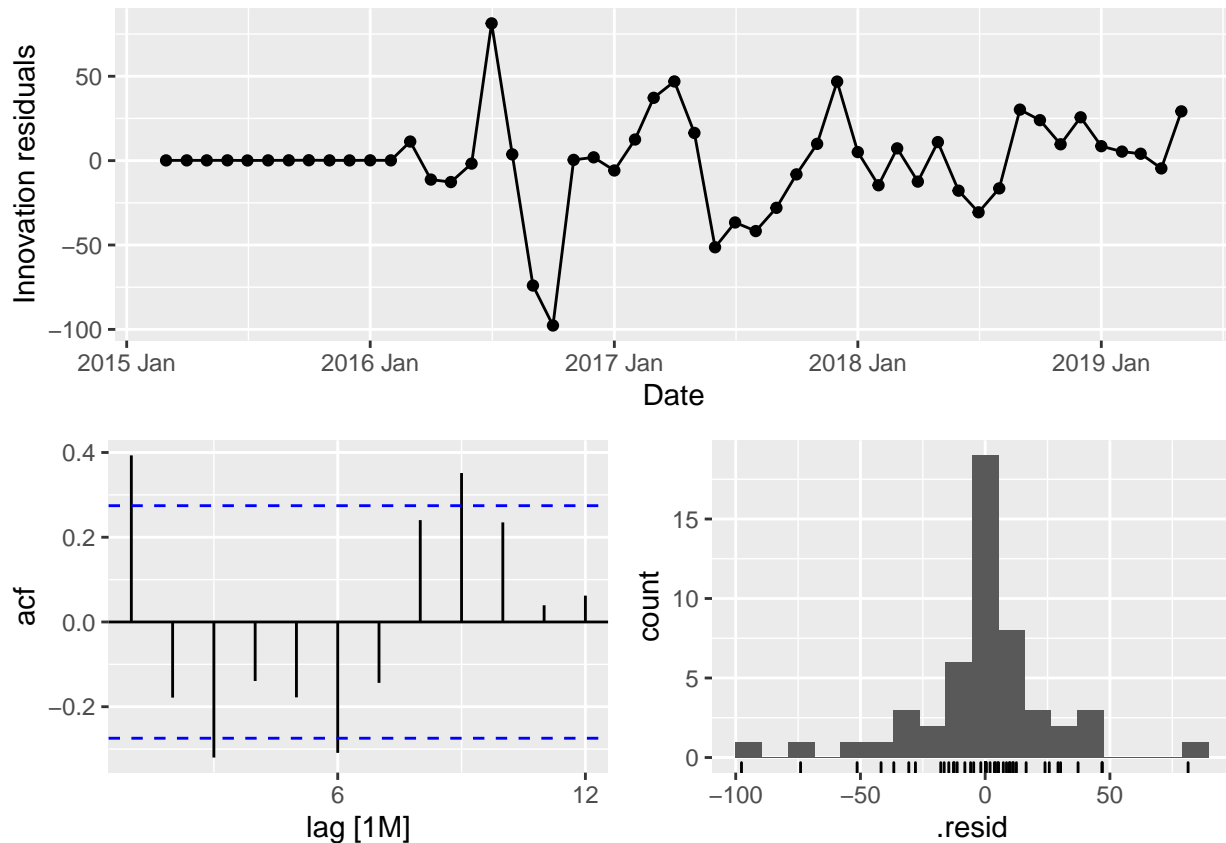
```
#AR and MA term
report(train |>
  filter(City == "Hyderabad") |>
  model(arima = ARIMA(AQI ~ pdq(0,0,0) + PDQ(1,1,1))))
```

```
## Series: AQI
## Model: ARIMA(0,0,0)(1,1,1)[12] w/ drift
##
## Coefficients:
##          sar1      sma1  constant
##         0.0623  -0.8626  -12.6512
## s.e.    0.5699   2.1957   3.3017
##
## sigma^2 estimated as 876.6:  log likelihood=-191.51
## AIC=391.01   AICc=392.19   BIC=397.67
```

We see that the seasonal terms considerably improved the AICc of the model. The model with only the seasonal MA term appears to be the best option given its lowest AICc value and being a simpler model. However, the seasonal MA term may not be significant given it is smaller than $2 \times \text{S.E.}$ The seasonal AR model does not seem to suffer from this

Let us check the residuals

```
gg_tsresiduals(train |>
  filter(City == "Hyderabad") |>
  model(arima = ARIMA(AQI ~ pdq(0,0,0) + PDQ(1,1,0))), lag=12)
```

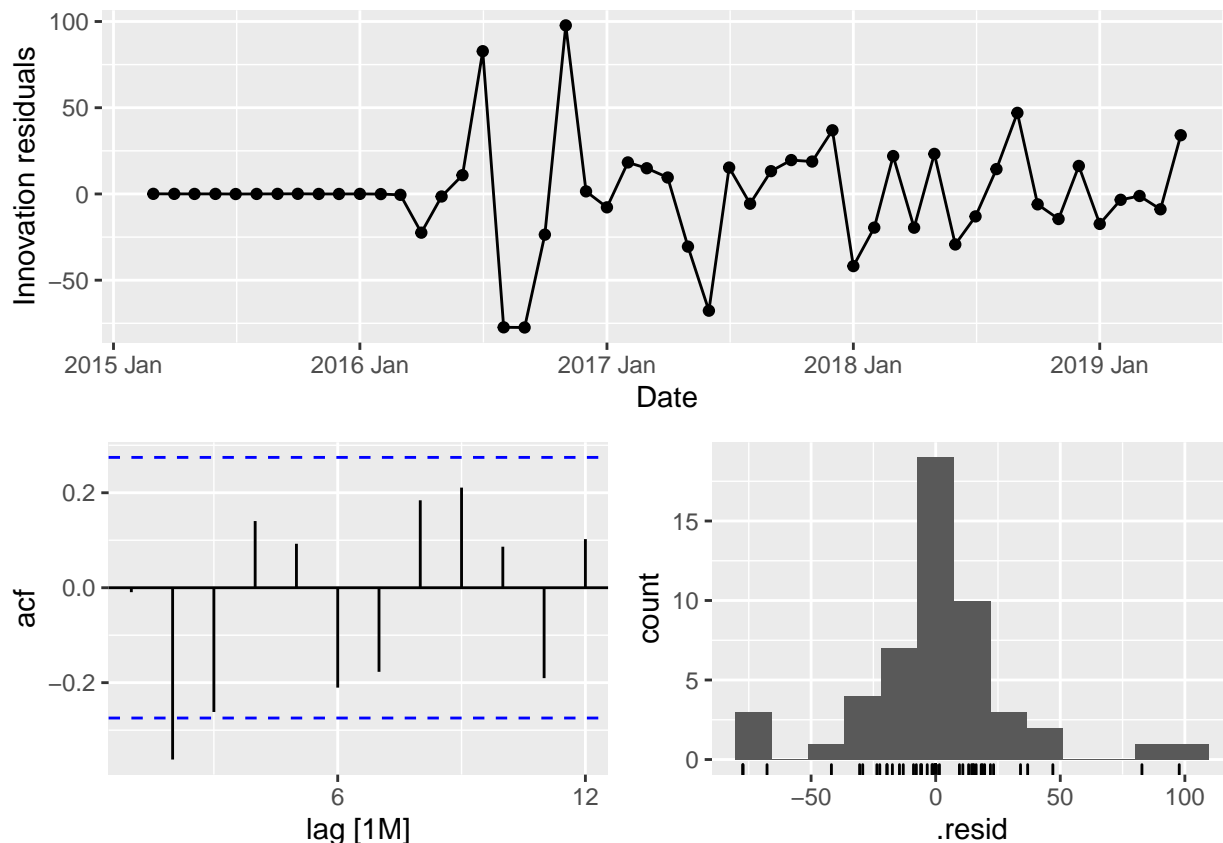


We see that there are some significant spikes in the residual ACF, spike 1 is more prominent. Perhaps we can apply a non-seasonal difference as well.

```
hyderabad.fit.seas <- train |>
  filter(City == "Hyderabad") |>
  model(arima = ARIMA(AQI ~ pdq(0,1,0) + PDQ(1,1,0)))

report(hyderabad.fit.seas)
```

```
## Series: AQI
## Model: ARIMA(0,1,0)(1,1,0)[12]
##
## Coefficients:
##      sar1
##    -0.4781
## s.e.   0.1619
##
## sigma^2 estimated as 1280: log likelihood=-190.91
## AIC=385.82  AICc=386.17  BIC=389.1
gg_tsresiduals(hyderabad.fit.seas, lag=12)
```



This model still sees a spike at lag 2. We could compare the forecasts of the seasonal and non-seasonal models for this city to see which one does better.

Using Automatic model Selection

```
fit <- train |> model(ARIMA(AQI, stepwise = FALSE))
fit |>
  pivot_longer(cols = -City, names_to = "Model name", values_to = "Model") |>
  mutate(model.details = format(Model))
```

```
## # A mable: 5 x 4
## # Key:      City, Model name [5]
##   City      `Model name`
##   <chr>     <chr>
## 1 Bengaluru ARIMA(AQI, stepwis~ <ARIMA(1,0,0)(1,0,0)[12] w/ mean> <ARIMA(1,0,0~
## 2 Chennai   ARIMA(AQI, stepwis~ <ARIMA(0,1,0)> <ARIMA(0,1,0~
## 3 Delhi     ARIMA(AQI, stepwis~ <ARIMA(0,0,1)(1,1,1)[12] w/ drift> <ARIMA(0,0,1~
## 4 Hyderabad ARIMA(AQI, stepwis~ <ARIMA(1,1,1)> <ARIMA(1,1,1~
## 5 Lucknow   ARIMA(AQI, stepwis~ <ARIMA(0,0,1)(0,1,1)[12]> <ARIMA(0,0,1~
```

```
glance(fit)
```

```
## # A tibble: 5 x 9
##   City      .model      sigma2 log_lik   AIC   AICc   BIC ar_roots ma_roots
##   <chr>     <chr>         <dbl>   <dbl> <dbl> <dbl> <dbl> <list>  <list>
## 1 Bengaluru ARIMA(AQI, stepw~    335.   -220.  449.  450.  456. <cpl>    <cpl>
## 2 Chennai   ARIMA(AQI, stepw~   1316.   -251.  503.  503.  505. <cpl>    <cpl>
## 3 Delhi     ARIMA(AQI, stepw~   1229.   -211.  431.  433.  440. <cpl>    <cpl>
```



```
## 4 Hyderabad ARIMA(AQI, stepw~ 745. -236. 478. 478. 483. <cpl> <cpl>
## 5 Lucknow ARIMA(AQI, stepw~ 1177. -195. 396. 396. 401. <cpl> <cpl>
```

Note that automatic selection did not pick a seasonal model for Hyderabad.


Additonal Analysis

We compare forecasts for Hyderabad using a seasonal and non-seasonal model

Hyderabad Non-seasonal vs Seasonal Model Forecast

```
hyderabad.fit |> forecast(new_data = test) |>
  autoplot(train)+
  geom_line(data = test |> filter(City == "Hyderabad"), aes(x=Date, y=AQI),
    color = "red", linetype = "longdash", alpha=0.7)
```

```
## `mutate_if()` ignored the following grouping variables:
## * Column `City`
```




```
hyderabad.fit |> forecast(new_data = test) |> accuracy(test)
```

```
## # A tibble: 1 x 11
##   .model City      .type      ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>   <chr>    <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 arima Hyderabad Test  -35.9  46.1  42.0 -65.4  69.9   NaN    NaN  0.644
```

```
hyderabad.fit.seas |> forecast(new_data = test) |>
  autoplot(train)+
```

```
geom_line(data = test |> filter(City == "Hyderabad"), aes(x=Date, y=AQI),
          color = "red", linetype = "longdash", alpha=0.7)
```

```
## `mutate_if()` ignored the following grouping variables:
## * Column `City`
```



```
hyderabad.fit.seas |> forecast(new_data = test) |> accuracy(test)
```

```
## # A tibble: 1 x 11
##   .model City      .type      ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>   <chr>    <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 arima  Hyderabad Test   -24.8  29.2  26.9 -37.3  38.9   NaN   NaN  0.0185
```

We see that the seasonal model does a much better job of forecasting the AQI than the non-seasonal model. We must note that the forecasts are off by a lot potentially because the test data is from the year 2020. We could try techniques other than ARIMA and compare their performance.


Forecasting the Remaining Series

We now conduct forecasts for the remaining cities and check model performance.

Delhi

```
delhi.fit |> forecast(new_data = test) |>
  autoplot(train)+
  geom_line(data = test |> filter(City == "Delhi"), aes(x=Date, y=AQI),
            color = "red", linetype = "longdash", alpha=0.7)
```

```
## `mutate_if()` ignored the following grouping variables:
## * Column `City`
```




```
delhi.fit |> forecast(new_data = test) |> accuracy(test)
```

```
## # A tibble: 1 x 11
##   .model City .type    ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>  <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 arima Delhi Test  -2.09  52.0  43.9 -4.15  29.4   NaN   NaN  0.534
```

Lucknow

```
lucknow.fit |> forecast(new_data = test) |>
  autoplot(train)+
  geom_line(data = test |> filter(City == "Lucknow"), aes(x=Date, y=AQI),
            color = "red", linetype = "longdash", alpha=0.7)
```

```
## `mutate_if()` ignored the following grouping variables:
## * Column `City`
```




```
lucknow.fit |> forecast(new_data = test) |> accuracy(test)
```

```
## # A tibble: 1 x 11
##   .model City   .type    ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>  <chr>  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 arima  Lucknow Test  -34.0  47.0  39.0 -23.4  27.5   NaN   NaN  0.646
```

Bengaluru

```
bengaluru.fit |> forecast(new_data = test) |>
  autoplot(train)+
  geom_line(data = test |> filter(City == "Bengaluru"), aes(x=Date, y=AQI),
    color = "red", linetype = "longdash", alpha=0.7)
```

```
## `mutate_if()` ignored the following grouping variables:
## * Column `City`
```




```
bengaluru.fit |> forecast(new_data = test) |> accuracy(test)
```

```
## # A tibble: 1 x 11
##   .model City      .type    ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>  <chr>    <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 arima  Bengaluru Test  -31.0  34.0  31.0 -48.0  48.0   NaN   NaN  0.434
```

Chennai

```
chennai.fit |> forecast(new_data = test) |>
  autoplot(train)+
  geom_line(data = test |> filter(City == "Chennai"), aes(x=Date, y=AQI),
    color = "red", linetype = "longdash", alpha=0.7)
```

```
## `mutate_if()` ignored the following grouping variables:
## * Column `City`
```



```
chennai.fit |> forecast(new_data = test) |> accuracy(test)
```

```
## # A tibble: 1 x 11
##   .model City   .type    ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>   <chr>   <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 arima  Chennai Test  -13.0  20.5  16.9 -17.9  21.2   NaN   NaN  0.335
```

Summary and Implications

We see that the series for Chennai was not suitably modeled by ARIMA. With no AR and MA terms it will be a flat line segment. We would expect the ARIMA models to work well with cities with well defined seasonal behavior such as Lucknow and Delhi, and even relatively less defined seasonal behavior for cities like Bengaluru and Hyderabad.

The forecasts show us that this is indeed true. Lucknow and Delhi have the best model performance as measured by accuracy metrics and the seasonal model is able to adequately capture the seasonality of the AQI (although the accuracy is not extremely high).

We see that the Bengaluru model performs poorly even though the seasonal pattern is captured by the model. We discuss the accuracy later in this section. Chennai's model has the highest accuracy but that is because the variation in AQI is not very high to begin with.

This analysis has shown us the importance of modeling seasonality. If we were to choose the automatic search model for Hyderabad or just use the ACF and PACF to decide which model to pick, we would have been left worse off with a model that would not be very reliable or helpful. The seasonality analysis helped us make a better model choice.

Lastly, given that the test data is from the year 2020, we should expect most models to perform poorly on this data. We see that this is true for all the series that we have forecasted. We can see that all the forecasts

for the year 2020 are much higher than the true values in the test data.

India had some of the most stringent lock down measures in place and the level of pollution dropped drastically in the early months of 2020 since economic activity mostly came to a standstill. Given that underlying data generating processes were dramatically altered, we must expect the performance of any forecasting model to be sub-optimal.

For next steps, we must test other forecasting techniques on this data. We could use machine learning based models, ETS or other techniques and check if they are able to improve accuracy. Another attempt to forecast could incorporate weather data - air quality is highly dependent on weather data. We could add other predictors with weather data to check if this helps improve model performance.

Given the forecasting accuracy and that the data is for the year 2020, these models would be useful if we had more years of data so that we could make longer forecasts and test long term accuracy of these models.